

Lecture 5:

Preprocessing and

sklearn pipelines

Varada Kolhatkar

Announcements

- HW1 grades have been posted.
- HW1 solutions have been posted on Canvas under Files tab. Please do not share them with anyone or do not post them anywhere.
- Syllabus quiz due date is September 19th, 11:59 pm.
- Homework 3 (hw3) has been released (Due: Oct 1st, 11:59 pm)
 - You can work in pairs for this assignment.

Recap

- Decision trees: Split data into subsets based on feature values to create decision rules
- k -NNs: Classify based on the majority vote from k nearest neighbors
- SVM RBFs: Create a boundary using an RBF kernel to separate classes

Recap

Aspect	Decision Trees	K-Nearest Neighbors (KNN)	Support Vector Machines (SVM) with RBF Kernel
Main hyperparameters	Max depth, min samples split	Number of neighbors (k)	C (regularization), Gamma (RBF kernel width)
Interpretability			
Handling of non- linearity			
Scalability			

Recap

Aspect	Decision Trees	K-Nearest Neighbors (KNN)	Support Vector Machines (SVM) with RBF Kernel
Sensitivity to outliers			
Memory usage			
Training time			
Prediction time			
Multiclass support			

(iClicker) Exercise 5.1

iClicker cloud join link: <https://join.iclicker.com/VYFJ>

Take a guess: In your machine learning project, how much time will you typically spend on data preparation and transformation?

- a. ~80% of the project time
- b. ~20% of the project time
- c. ~50% of the project time
- d. None. Most of the time will be spent on model building

The question is adapted from [here](#).

(iClicker) Exercise 5.2

iClicker cloud join link: <https://join.iclicker.com/VYFJ>

Select all of the following statements which are TRUE.

- a. **StandardScaler** ensures a fixed range (i.e., minimum and maximum values) for the features.
- b. **StandardScaler** calculates mean and standard deviation for each feature separately.
- c. In general, it's a good idea to apply scaling on numeric features before training k-NN or SVM RBF models.
- d. The transformed feature values might be hard to interpret for humans.
- e. After applying **SimpleImputer** The transformed data has a different shape than the original data.

(iClicker) Exercise 5.3

iClicker cloud join link: <https://join.iclicker.com/VYFJ>

Select all of the following statements which are TRUE.

- a. You can have scaling of numeric features, one-hot encoding of categorical features, and scikit-learn estimator within a single pipeline.
- b. Once you have a `scikit-learn` pipeline object with an estimator as the last step, you can call `fit`, `predict`, and `score` on it.
- c. You can carry out data splitting within `scikit-learn` pipeline.
- d. We have to be careful of the order we put each transformation and model in a pipeline.

Preprocessing motivation: example

You're trying to find a suitable date based on:

- Age (closer to yours is better).
- Number of Facebook Friends (closer to your social circle is ideal).

Preprocessing motivation: example

- You are 30 years old and have 250 Facebook friends.

Person	Age	#FB Friends	Euclidean Distance Calculation	Distance
A	25	400	$\sqrt{(5^2 + 150^2)}$	150.08
B	27	300	$\sqrt{(3^2 + 50^2)}$	50.09
C	30	500	$\sqrt{(0^2 + 250^2)}$	250.00
D	60	250	$\sqrt{(30^2 + 0^2)}$	30.00

Based on the distances, the two nearest neighbors (2-NN) are:

- Person D** (Distance: 30.00)
- Person B** (Distance: 50.09)

What's the problem here?

Common transformations

Imputation: Fill the gaps! ()

Fill in missing data using a chosen strategy:

- **Mean:** Replace missing values with the average of the available data.
- **Median:** Use the middle value.
- **Most Frequent:** Use the most common value (mode).
- **KNN Imputation:** Fill based on similar neighbors.

Example:

Imputation is like filling in your average or median or most frequent grade for an assessment you missed.

```
1 from sklearn.impute import SimpleImputer
2 imputer = SimpleImputer(strategy='mean')
3 X_imputed = imputer.fit_transform(X)
```

Scaling: Everything to the same range! ()

Ensure all features have a comparable range.

- **StandardScaler**: Mean = 0, Standard Deviation = 1.

Example:

Scaling is like adjusting the number of everyone's Facebook friends so that both the number of friends and their age are on a comparable scale. This way, one feature doesn't dominate the other when making comparisons.

```
1 from sklearn.preprocessing import StandardScaler
2 scaler = StandardScaler()
3 X_scaled = scaler.fit_transform(X)
```

One-Hot encoding: 🍏 →

Convert categorical features into binary columns.

- Creates new binary columns for each category.
- Useful for handling categorical data in machine learning models.

Example:

Turn “Apple, Banana, Orange” into binary columns:

Fruit	🍏	🍌	🍊
Apple 🍏	1	0	0
Banana 🍌	0	1	0
Orange 🍊	0	0	1

```
1 from sklearn.preprocessing import OneHotEncoder
2 encoder = OneHotEncoder()
3 X_encoded = encoder.fit_transform(X)
```

Ordinal encoding: Ranking matters!



Convert categories into integer values that have a meaningful order.

- Assign integers based on order or rank.
- Useful when there is an inherent ranking in the data.

Example:

Turn “Poor, Average, Good” into 1, 2, 3:

Rating	Ordinal
Poor	1
Average	2
Good	3

```
1 from sklearn.preprocessing import OrdinalEncoder
2 encoder = OrdinalEncoder()
3 X_ordinal = encoder.fit_transform(X)
```

sklearn Transformers vs Estimators

Transformers

- Are used to transform or preprocess data.
- Implement the `fit` and `transform` methods.
 - `fit(X)`: Learns parameters from the data.
 - `transform(X)`: Applies the learned transformation to the data.
- Examples:
 - Imputation (`SimpleImputer`): Fills missing values.
 - Scaling (`StandardScaler`): Standardizes features.

Estimators

- Used to make predictions.
- Implement `fit` and `predict` methods.
 - `fit(X, y)`: Learns from labeled data.
 - `predict(X)`: Makes predictions on new data.
- Examples: `DecisionTreeClassifier`, `SVC`, `KNeighborsClassifier`

The golden rule in feature transformations

- Never transform the entire dataset at once!
- **Why?** It leads to **data leakage** — using information from the test set in your training process, which can artificially inflate model performance.
- **Fit** transformers like scalers and imputers on the **training set only**.
- **Apply** the transformations to both the training and test sets **separately**.

Example:

```
1 from sklearn.preprocessing import StandardScaler
2 scaler = StandardScaler()
3 X_train_scaled = scaler.fit_transform(X_train)
4 X_test_scaled = scaler.transform(X_test)
```

sklearn Pipelines

- Pipeline is a way to chain multiple steps (e.g., preprocessing + model fitting) into a single workflow.
- Simplify the code and improves readability.
- Reduce the risk of data leakage by ensuring proper transformation of the training and test sets.
- Automatically apply transformations in sequence.

Example:

Chaining a `StandardScaler` with a `KNeighborsClassifier` model.

```
1 from sklearn.pipeline import make_pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.neighbors import KNeighborsClassifier
4
5 pipeline = make_pipeline(StandardScaler(), KNeighborsClassifier())
6
7 pipeline.fit(X_train, y_train)
8 y_pred = pipeline.predict(X_test)
```

Class demo