

CPSC 330 Lecture 8: Hyperparameter Optimization

Varada Kolhatkar

Announcements

- Important information about midterm 1
 - <https://piazza.com/class/m01ukubppof625/post/249>
- Change of my office hours
 - Thursdays from 2 to 3 in my office ICCS 237
- HW3 is due today 11:59 pm.
- HW4 has been released

Recap: Logistic regression

- A linear model used for binary classification tasks.
 - (Optional) There is an extension of logistic regression called multinomial logistic regression for multiclass classification.
- Parameters:
 - **Coefficients (Weights):** The model learns a coefficient or a weight associated with each feature that represents its importance.
 - **Bias (Intercept):** A constant term added to the linear combination of features and their coefficients.

Recap: Logistic regression

- The model computes a weighted sum of the input features' values, adjusted by their respective coefficients and the bias term.
- This weighted sum is passed through a sigmoid function to transform it into a probability score, indicating the likelihood of the input belonging to the “positive” class.

$$P_{\text{hat}} = \sigma \left(\sum_{i=1}^d w_i x_i + b \right)$$

- P_{hat} is the predicted probability of the example belonging to the positive class.
- w_i is the learned weight associated with feature i
- x_i is the value of the input feature i
- b is the bias term

Recap: Logistic regression

- For a dataset with d features, the decision boundary that separates the classes is a $d - 1$ dimensional hyperplane.
- Complexity hyperparameter: C in `sklearn`.
 - Higher $C \rightarrow$ more complex model meaning larger coefficients
 - Lower $C \rightarrow$ less complex model meaning smaller coefficients

Recap: **CountVectorizer** input

- Primarily designed to accept either a **pandas.Series** of text data or a 1D **numpy** array. It can also process a list of string data directly.
- Unlike many transformers that handle multiple features (**DataFrame** or 2D **numpy** array), **CountVectorizer** a single text column at a time.
- If your dataset contains multiple text columns, you will need to instantiate separate **CountVectorizer** objects for each text feature.
- This approach ensures that the unique vocabulary and tokenization processes are correctly applied to each specific text column without interference.

Hyperparameter optimization motivation

Data

```

1 sms_df = pd.read_csv(DATA_DIR + "spam.csv", encoding="latin-1")
2 sms_df = sms_df.drop(columns = ["Unnamed: 2", "Unnamed: 3", "Unnamed: 4"])
3 sms_df = sms_df.rename(columns={"v1": "target", "v2": "sms"})
4 train_df, test_df = train_test_split(sms_df, test_size=0.10, random_state=42)
5 X_train, y_train = train_df["sms"], train_df["target"]
6 X_test, y_test = test_df["sms"], test_df["target"]
7 train_df.head(4)

```

	target	sms
3130	spam	LookAtMe!: Thanks for your purchase of a video...
106	ham	Aight, I'll hit you up when I get some cash
4697	ham	Don no da:)whats you plan?
856	ham	Going to take your babe out ?

Model building

- Let's define a pipeline

```
1 pipe_svm = make_pipeline(CountVectorizer(), SVC())
```

- Suppose we want to try out different hyperparameter values.

```
1 parameters = {  
2     "max_features": [100, 200, 400],  
3     "gamma": [0.01, 0.1, 1.0],  
4     "C": [0.01, 0.1, 1.0],  
5 }
```

Hyperparameter optimization with loops

- Define a parameter space.
- Iterate through possible combinations.
- Evaluate model performance.
- What are some limitations of this approach?

sklearn methods

- `sklearn` provides two main methods for hyperparameter optimization
 - Grid Search
 - Random Search

Grid Search

- Covers all possible combinations from the provided grid.
- Can be parallelized easily.
- Integrates cross-validation.

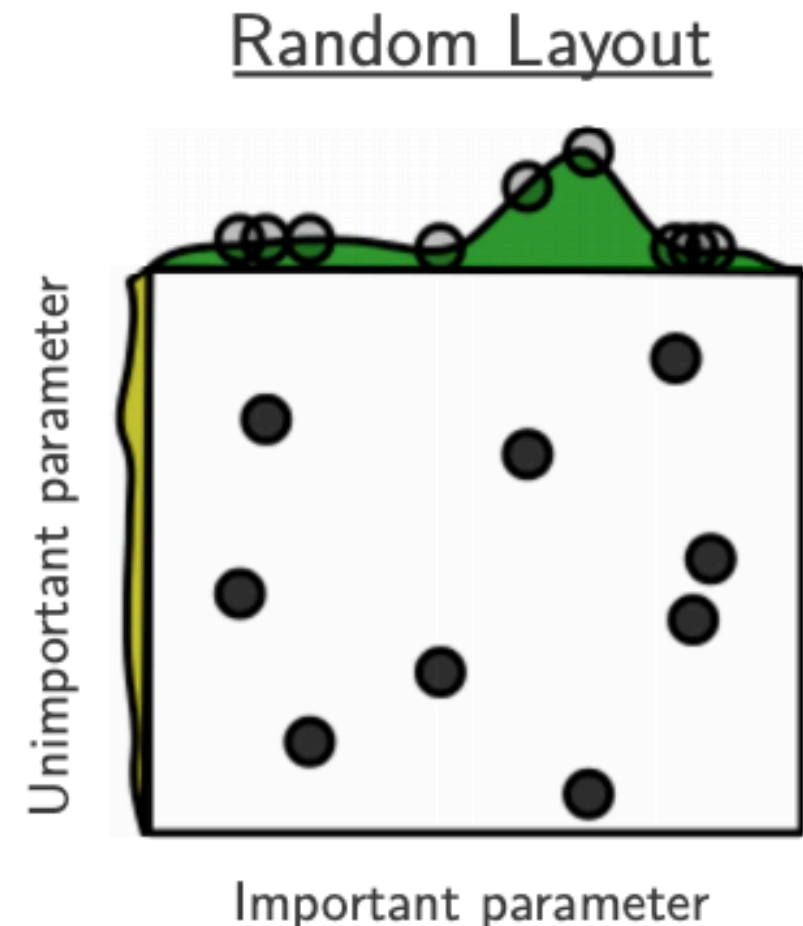
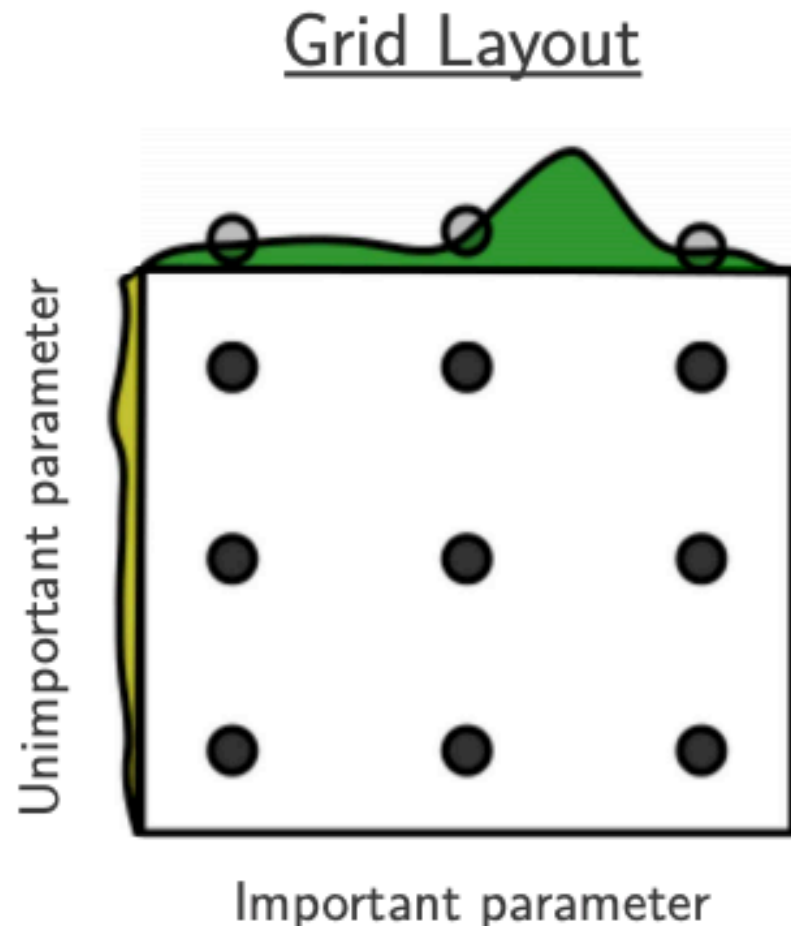
Grid search example

```
1 from sklearn.model_selection import GridSearchCV
2
3 pipe_svm = make_pipeline(CountVectorizer(), SVC())
4
5 param_grid = {
6     "countvectorizer__max_features": [100, 200, 400],
7     "svc__gamma": [0.01, 0.1, 1.0],
8     "svc__C": [0.01, 0.1, 1.0],
9 }
10 grid_search = GridSearchCV(pipe_svm,
11                             param_grid = param_grid,
12                             n_jobs=-1,
13                             return_train_score=True
14                             )
15 grid_search.fit(X_train, y_train)
16 grid_search.best_score_
```

0.9782606272997375

Random Search

- More efficient than grid search when dealing with large hyperparameter spaces.
- Samples a given number of parameter settings from distributions.



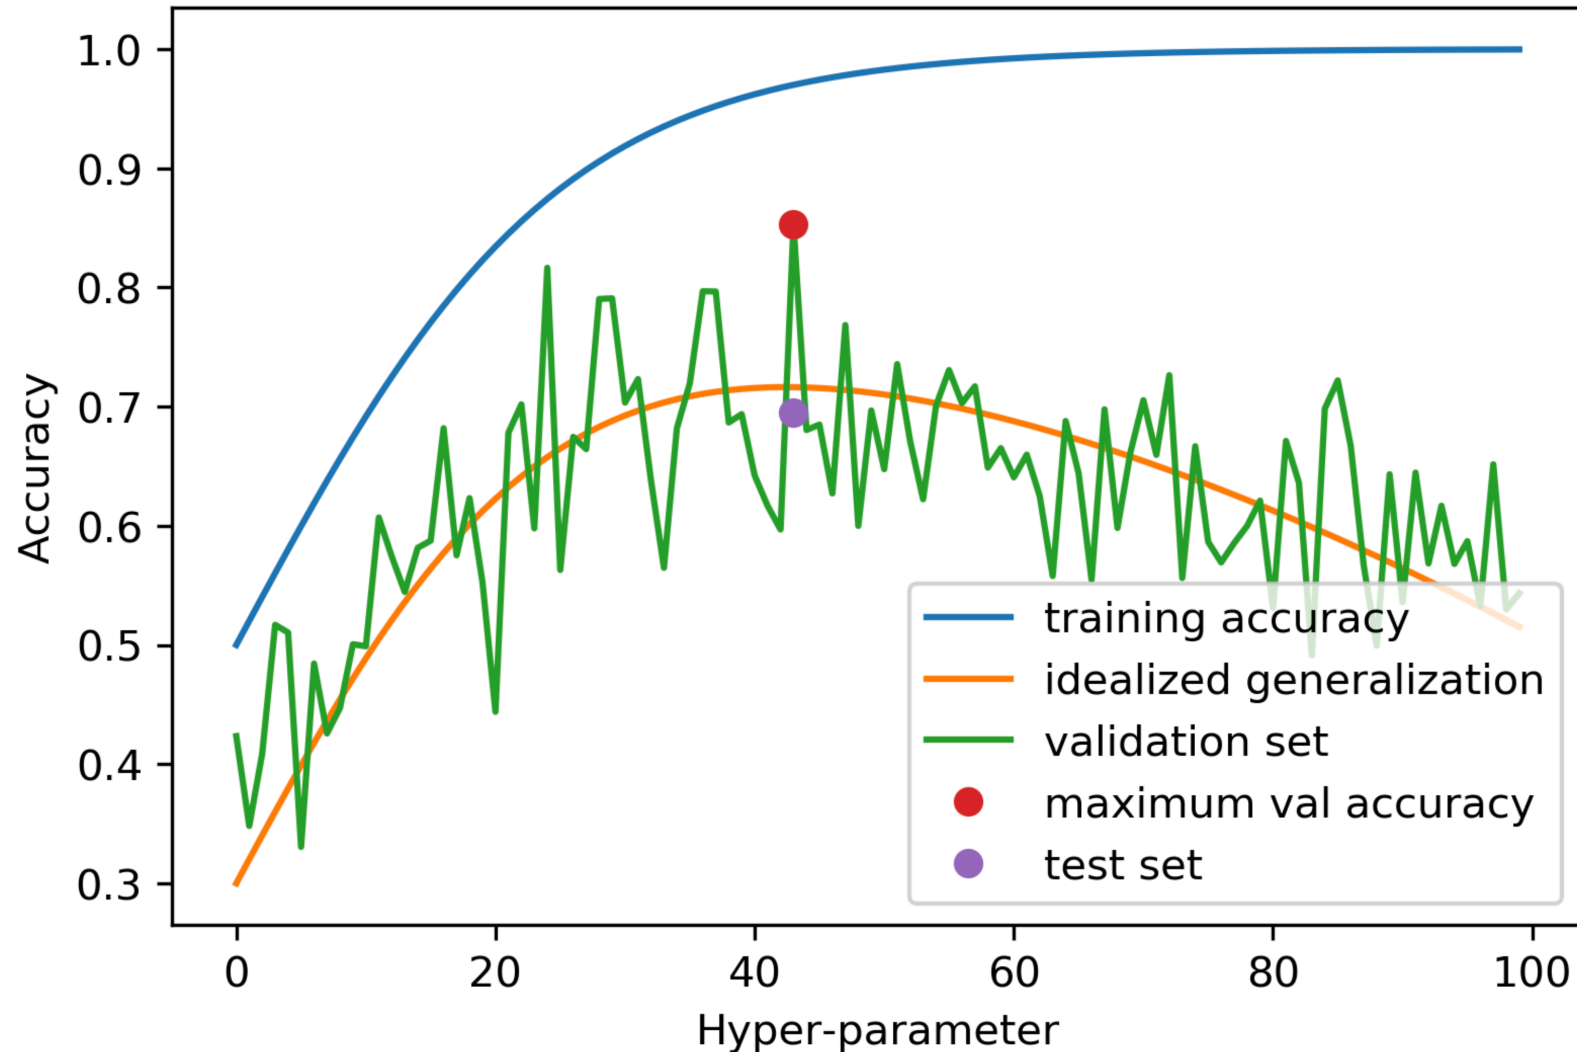
Random search example

```
1 from sklearn.model_selection import RandomizedSearchCV
2 from scipy.stats import randint
3
4 pipe_svc = make_pipeline(CountVectorizer(), SVC())
5
6 param_dist = {
7     "countvectorizer__max_features": randint(100, 2000),
8     "svc__C": uniform(0.1, 1e4), # loguniform(1e-3, 1e3),
9     "svc__gamma": loguniform(1e-5, 1e3),
10 }
11 random_search = RandomizedSearchCV(pipe_svc,
12                                   param_distributions = param_dist,
13                                   n_iter=10,
14                                   n_jobs=-1,
15                                   return_train_score=True)
16
17 # Carry out the search
18 random_search.fit(X_train, y_train)
19 random_search.best_score_
```

0.9834462679824798

Optimization bias

- Why do we need separate validation and test datasets?



Mitigating optimization bias.

- Cross-validation
- Ensembles
- Regularization and choosing a simpler model

(iClicker) Exercise 8.1

iClicker cloud join link: <https://join.iclicker.com/VYFJ>

Select all of the following statements which are TRUE.

- a. If you get best results at the edges of your parameter grid, it might be a good idea to adjust the range of values in your parameter grid.
- b. Grid search is guaranteed to find the best hyperparameter values.
- c. It is possible to get different hyperparameters in different runs of RandomizedSearchCV.

Questions for you

- You have a dataset and you give me 1/10th of it. The dataset given to me is rather small and so I split it into 96% train and 4% validation split. I carry out hyperparameter optimization using a single 4% validation split and report validation accuracy of 0.97. Would it classify the rest of the data with similar accuracy?
 - Probably
 - Probably not

Questions for class discussion

- Suppose you have 10 hyperparameters, each with 4 possible values. If you run `GridSearchCV` with this parameter grid, how many cross-validation experiments will be carried out?
- Suppose you have 10 hyperparameters and each takes 4 values. If you run `RandomizedSearchCV` with this parameter grid with `n_iter=20`, how many cross-validation experiments will be carried out?

Class Demo